

libCEED: A Case Study in Hidden Benefits of the Bridge Pattern

Jeremy L Thompson,
Natalie Beams, Jed Brown, and Yohann Dudouit

University of Colorado Boulder

jeremy.thompson@colorado.edu

July 31, 2020

libCEED Team

Developers: Ahmad Abdelfattah¹, Valeria Barra²,
Natalie Beams¹, Jed Brown², Jean-Sylvain Camier³,
Veselin Dobrev³, Yohann Dudouit³, Leila Ghaffari²,
Tzanio Kolev³, David Medina⁴, Thilina Rathnayake⁵,
Jeremy L. Thompson², & Stan Tomov⁵

Grant: Exascale Computing Project (17-SC-20-SC)

- 1: University of Tennessee
- 2: University of Colorado, Boulder
- 3: Lawrence Livermore National Laboratory
- 4: Occalytics LLC
- 5: University of Illinois at Urbana-Champaign

libCEED is an extensible library that provides a portable algebraic interface and optimized implementations of high-order operators

We have optimized implementations for multiple architectures

Bridge design pattern offers performance portability
as well as improved debugability and internal design documentation

Overview

- 1 Introduction
- 2 libCEED Design
- 3 Backend Development
- 4 Future Work
- 5 Questions

Center for Efficient Exascale Discretizations

DoE exascale co-design center

- Design discretization algorithms for exascale hardware that deliver significant performance gain over low order methods
- Collaborate with hardware vendors and software projects for exascale hardware and software stack
- Provide efficient and user-friendly unstructured PDE discretization component for exascale software ecosystem



libCEED Philosophy

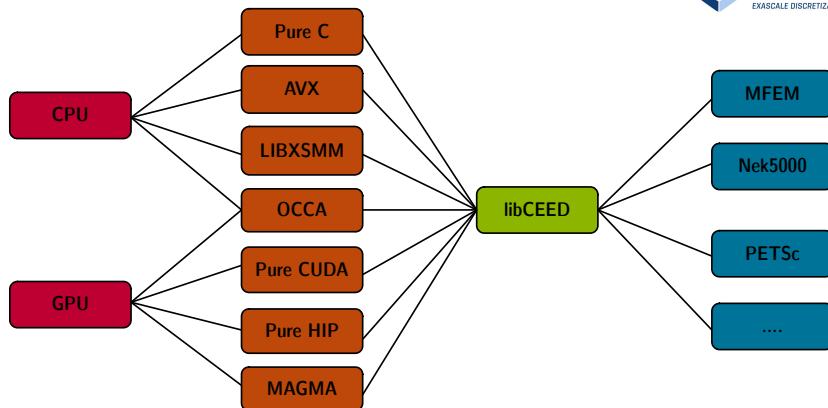
libCEED provides purely algebraic interface for matrix-free evaluation of arbitrary polynomial order PDE operators

libCEED design approach:

- Optimized implementations for multiple architectures
- Runtime selection of backend implementation
- Single source user PDE quadrature point functions

Repository: <https://github.com/CEED/libCEED>

libCEED Backends




Natural use case for bridge pattern

libCEED Operator Decomposition



$$A = P^T G^T B^T D B G P$$



$$A_L = G^T B^T D B G$$

- G - CeedElemRestriction, local gather/scatter
- B - CeedBasis, provides basis operations such as interp and grad
- D - CeedQFunction, representation of PDE at quadrature points
- A_L - CeedOperator, aggregation of Ceed objects for local action of operator

User QFunction Source

User QFunction Code:

```
// ---- Fuvisc
const CeedInt Fuviscidx[3][3] = {{0, 1, 2}, {1, 3, 4}, {2, 4, 5}};
for (CeedInt j=0; j<3; j++)
  for (CeedInt k=0; k<3; k++)
    dv[k][j+1][i] -= wJ*(Fu[Fuviscidx[j][0]]*dXdxT[k][0] +
                        Fu[Fuviscidx[j][1]]*dXdxT[k][1] +
                        Fu[Fuviscidx[j][2]]*dXdxT[k][2]);
```

Compiled Assembly:

```
dv[k][j+1][i] -= wJ*(Fu[Fuviscidx[j][0]]*dXdxT[k][0] +
b08d:  c5 7d 28 d0          vmovapd %ymm0,%ymm10
                        Fu[Fuviscidx[j][1]]*dXdxT[k][1] +
b091:  c4 42 c5 b8 d3      vfmadd231pd %ymm11,%ymm7,%ymm10
b096:  c5 fd 28 84 24 c8 04  vmovapd 0x4c8(%rsp),%ymm0
b09d:  00 00
dv[k][j+1][i] -= wJ*(Fu[Fuviscidx[j][0]]*dXdxT[k][0] +
b09f:  c4 62 f5 ac 14 07      vfnmadd213pd (%rdi,%rax,1),%ymm1,%ymm10
b0a5:  c5 7d 11 14 07        vmovupd %ymm10,(%rdi,%rax,1)
                        Fu[Fuviscidx[j][1]]*dXdxT[k][1] +
b0aa:  c5 7d 59 94 24 68 04  vmulpd 0x468(%rsp),%ymm0,%ymm10
b0b1:  00 00
```

Similar optimized code on other architectures

End Goal: Optimized Backends

CPU:

- AVX instructions for vector operations, FMA
- LIBXSMM for JiT optimized small matrix multiplication

GPU:

- JiT to generate device kernels from user code
- Fused kernels to minimize launches and data movement

Development Challenges

CPU:

- Optimized code can be difficult to read
- External library calls can be opaque

GPU:

- Debugging more challenging on devices
- Fused kernels can't be incrementally tested

Bridge Pattern FTW

Strong encapsulation of backend implementation details from interface

- Reference backends provide straightforward implementation
 - /cpu/self/ref/serial
 - /gpu/cuda/ref
 - /gpu/hip/ref
- Debugging backends provide additional debugging tools
 - /cpu/self/memcheck
- Optimized backends selectively re-implement objects
 - /cpu/self/xsmm/blocked
 - /gpu/cuda/gen

Hidden Benefits

- Stable API and ABI across implementations
- Flexible and distributed development
- Clear reference code for new developers
- Improved debugability ,for users and developers
- ...

Future Work

- Further performance enhancements (GPU and CPU)
- Improved mixed mesh and operator composition support
- Expanded non-linear and multi-physics examples
- Preconditioning based on libCEED operator decomposition
- Algorithmic differentiation of user quadrature functions
- We invite contributors and friendly users

Questions?

Advisors : Jed Brown¹ & Adriana Gillman¹

Developers: Ahmad Abdelfattah¹, Valeria Barra²,
Natalie Beams¹, Jed Brown², Jean-Sylvain Camier³,
Veselin Dobrev³, Yohann Dudouit³, Leila Ghaffari²,
Tzanio Kolev³, David Medina⁴, Thilina Rathnayake⁵,
& Stan Tomov⁵

Grant: Exascale Computing Project (17-SC-20-SC)

- 1: University of Tennessee
- 2: University of Colorado, Boulder
- 3: Lawrence Livermore National Laboratory
- 4: Occalytics LLC
- 5: University of Illinois at Urbana-Champaign

libCEED: A Case Study in Hidden Benefits of the Bridge Pattern

Jeremy L Thompson,
Natalie Beams, Jed Brown, and Yohann Dudouit

University of Colorado Boulder

jeremy.thompson@colorado.edu

July 31, 2020