

# On Performance and Portability for Generic Finite Element Interfaces

Jeremy L Thompson

University of Colorado Boulder

*[jeremy.thompson@colorado.edu](mailto:jeremy.thompson@colorado.edu)*

July 11, 2018

A global sparse matrix is no longer a good representation of a high-order linear operator

libCEED is an extensible library that provides a portable algebraic interface and optimized implementations

We have preliminary results comparing performance to native implementations in production software

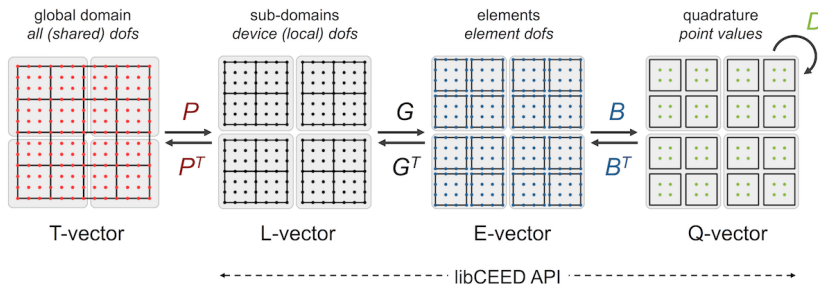
# Overview

- 1 Introduction
- 2 libCEED
- 3 Production Software
- 4 Performance Comparison
- 5 Questions

# Operator Decomposition




$$A = P^T G^T B^T D B G P$$



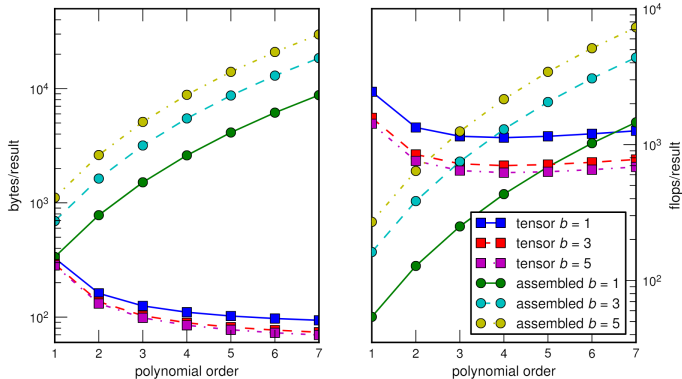
# Matrix Free Implementation

$$A = P^T G^T B^T D B G P$$



- Avoid global matrix assembly
- Map each element to reference element
- Only store map to reference, action on reference
- Easy to parallelize across nodes

# Assembled Matrix Cost!



Memory bandwidth and ops per dof to apply a Jacobian from  $Q_k$  discretization of a  $b$ -variable PDE system using an assembled matrix versus matrix-free exploiting the tensor product structure

# libCEED API

- Provides on-device operator implementation
- Easy to incorporate into existing code
- Supports multiple types of computational devices
  - CPU - Reference and vectorized, template for new backends
  - OCCA (jit) - CPU, OpenMP, OpenCL, and CUDA
  - MAGMA
  - One source code can call multiple CEEDs with different backends
- v0.2 March and v0.3 (imminent)
- BSD-2 license

# API Objects

- $G$  - CeedRestriction  
Restrict to single element
- $B$  - CeedBasis  
Actions on basis such as interpolation,  
gradient, divergence, curl
- $D$  - CeedQFunction  
Operator action at quadrature points  
to include coefficient functions



# Device Level Operator

- $L = G^T B^T D B G$  - CeedOperator
- libCEED objects are combined to create a CeedOperator
- CeedOperator gives operator action for elements on device
- User code responsible for communication between devices

$$A = P^T L P$$

# Basis

- Tensor  $H^1$  elements
- User provides  $p, q, dim$  and chooses Gauss or Gauss-Lobatto dofs
- Alternatively, user provides  $1D$  interp, grad matrices and quadrature weights and points
- More geometries,  $H(div)$ ,  $H(curl)$  coming

```
CeedBasisCreateTensorH1Lagrange(ceed, dim, ncomp,
                                P, Q, CEED_GAUSS, &basis);
CeedBaisApply(basis, CEED_NOTRANSPOSE, CEED_EVAL_INTERP,
              x, xq);
```

# Restriction

- Gather and scatter operations
- Support conforming and non-conforming meshes
- User provides index list, may be linear combination of dofs
- On node communication only

```
CeedElemRestrictionCreate(ceed, ne, 2, ne+1, 1,  
                          CEED_MEM_HOST, CEED_USE_POINTER,  
                          ind, &rstr);  
CeedElemRestrictionApply(rstr, CEED_NOTRANSPOSE,  
                          CEED_NOTRANSPOSE, u, ru,  
                          CEED_REQUEST_IMMEDIATE);
```

# Qfunction

- Applies the physics at the quadrature points
- Multiple inputs and outputs

```
CeedQFunctionCreateInterior(ceed, 1, myfunc,  
                           __FILE__ ":myfunc", &qf);  
CeedQFunctionAddInput(qf_setup, "_weight",  
                      1, CEED_EVAL_WEIGHT);  
CeedQFunctionAddInput(qf_setup, "x",  
                      1, CEED_EVAL_GRAD);  
CeedQFunctionAddOutput(qf_setup, "rho",  
                       1, CEED_EVAL_NONE);
```

# Operator

- Combines components to give local operator action
- Multiple inputs and outputs
- Composite operators coming

```
CeedOperatorCreate(ceed, qf_setup, NULL, NULL, &op_setup);
```

```
CeedOperatorSetField(op_setup, "_weight",  
                     CEED_RESTRICTION_IDENTITY, basis,  
                     CEED_VECTOR_NONE);
```

```
CeedOperatorSetField(op_setup, "x",  
                     rstr, basis, CEED_VECTOR_ACTIVE);
```

```
CeedOperatorSetField(op_setup, "rho",  
                     CEED_RESTRICTION_IDENTITY,  
                     CEED_BASIS_COLOCATED,  
                     CEED_VECTOR_ACTIVE);
```

# Benefits

- Extensible library
- Lower memory transfer, no sparse matrix
- Implementations for multiple devices and backends
- Backend improvements benefit all applications
  - Tensor contraction, basis application, etc
- Minimal dependencies

# Standalone Implementation

```
// Create the mass operator.  
CeedOperator oper;  
CeedOperatorCreate(CEED, apply_qfunc ,  
                  NULL, NULL, &oper);  
  
...  
  
// Apply the mass operator: 'u'  $\rightarrow$  'v'.  
CeedOperatorApply(oper, u, v,  
                  CEED_REQUEST_IMMEDIATE);
```

# MFEM

```
/// Wrapper for a mass CeedOperator as an  
/// mfem::Operator  
class CeedMassOperator : public mfem::Operator  
protected:  
    const mfem::FiniteElementSpace *fes;  
    CeedOperator build_oper, oper;  
    CeedBasis basis, mesh_basis;  
    CeedElemRestriction restr, mesh_restr;  
    CeedQFunction apply_qfunc, build_qfunc;  
    CeedVector node_coords, qdata;
```



## Nek5000

```
subroutine ceed_axhm1(pap,ap1,p1,h1,h2,ceed,op_mass,
$ vec_ap1,vec_p1,vec_qdata)
```

```
include 'ceedf.h'
```

```
c Vector conjugate gradient matvec for solution of
c uncoupled Helmholtz equations
```

```
include 'SIZE'
```

```
include 'TOTAL'
```

```
...
```

```
call ceedvectorsetarray(vec_p1,ceed_mem_host,
$ ceed_use_pointer, p1,err)
```

```
call ceedoperatorapply(op_mass,vec_p1,vec_ap1,
$ ceed_request_immediate,err)
```

```
call ceedvectorgetarray(vec_ap1,ceed_mem_host,ap1,err)
```

# PETSc

```

user->op = op_mass;
user->qdata = qdata;

ierr = MatCreateShell(comm, mdof[0]*mdof[1]*mdof[2],
                      mdof[0]*mdof[1]*mdof[2],
                      PETSC_DECIDE, PETSC_DECIDE, user, &mat);
CHKERRQ(ierr);
ierr = MatShellSetOperation(mat, MATOP_MULT
                           (void*)(void)) MatMult_Mass); CHKERRQ(ierr);

...

ierr = KSPSetFromOptions(ksp); CHKERRQ(ierr);
ierr = KSPSetOperators(ksp, mat, mat); CHKERRQ(ierr);
ierr = KSPSolve(ksp, rhs, X); CHKERRQ(ierr);

```

# Nek5000



Research Computing  
UNIVERSITY OF COLORADO **BOULDER**



Problem:  $\nabla u = f$   
CEED Benchmark Problem 1

Computer: CU Boulder Summit

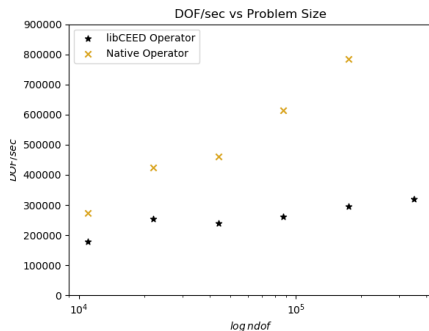

Domain: 3D Cube  
Elements: Hexagonal  
Number of Elements:  $2^n$   
Shape Function Order: 7  
Quadrature Points: 9

Nodes: 1  
CPUs: Intel Xeon "Haswell"  
Processors: 32  
Compiler: Intel/17.0.0  
MPI: Intel/2017.0.098

# Nek5000 - The Good News



# Nek5000 - The Bad News



# Future Work

- Improve optimized CPU backend, vectorize across elements
- Improve GPU backends, reduce data movement
- Add additional geometries, tets, pyramids, and prisms
- Create library of user quadrature functions
- Composite operators, for mixed meshes and multiphysics
- Create pure CUDA backend
- Compare libCEED operators to native implementation in a wider range of production software
- Contributors and friendly users welcome

# Questions?

Advisor: Jed Brown<sup>1</sup>

Collaborators: Jean-Sylvain Camier<sup>2</sup>, Tzanio Kolev<sup>2</sup>,  
Veselin Dobrev<sup>2</sup>, & Thilina Rathnayake<sup>3</sup>

Grant: Exascale Computing Project (17-SC-20-SC)

1: University of Colorado, Boulder

2: Lawrence Livermore National Laboratory

3: University of Illinois, Urbana-Champaign

# On Performance and Portability for Generic Finite Element Interfaces

Jeremy L Thompson

University of Colorado Boulder

*[jeremy.thompson@colorado.edu](mailto:jeremy.thompson@colorado.edu)*

July 11, 2018